

1. What is <Name>

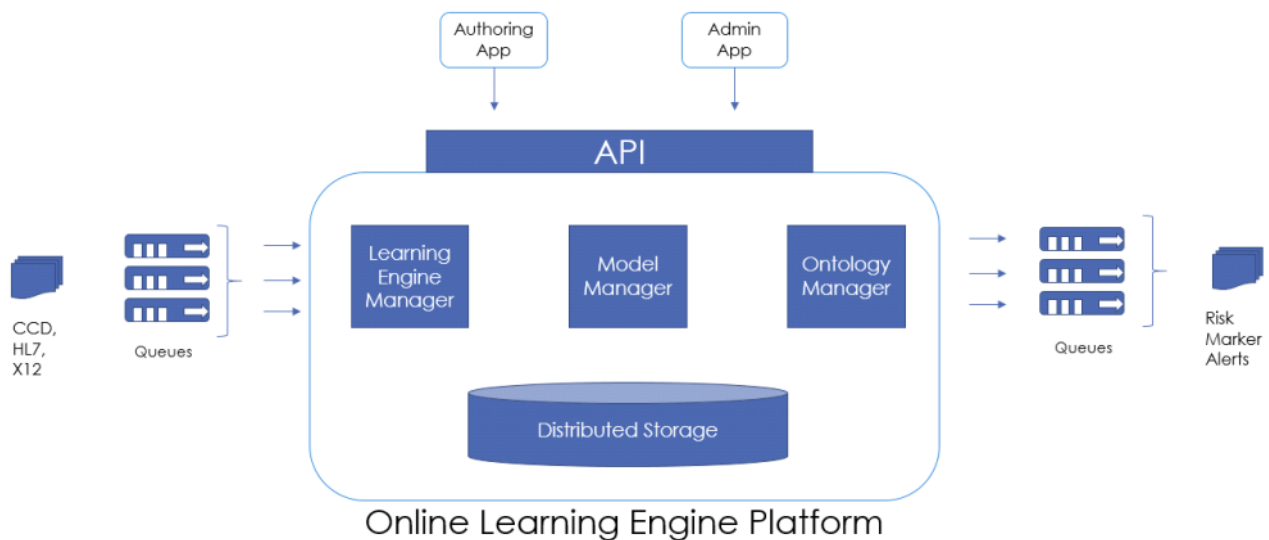
Friday, October 31, 2014 11:20 AM

Overview

<Name> is an online learning engine platform that provides near real-time processing streaming decision management, and produces actionable or programmable alerts based on your business rule sets. It accepts <XCDR, CCD, HL7, X12 *other xml?*> message inputs, and evaluates them against models based on pre-defined rule sets and historical data. <Name> generates and publishes real-time alerts based on the outcome of these models. These alerts are pushed into a messaging queue which can be consumed via downstream applications such as SMS portals, mobile applications, dashboards, business action solutions or automated expert action systems.

It has been used to help one financial customer reduce overnight global system transaction analysis from over 10 minutes to under 5, resulting in millions saved each period. It is also being used in medical applications to create alerts for possible undiagnosed conditions.

<Product Name> Architecture



Q: Is our Decision Engine chainable? Need language that clarifies how an Decision engine is different than a Decision Tree?

Q: What is the plan for our User Interface for query design? ([Component #3](#))

Q: Is <Name> stand-alone and embeddable? ([JasperReports Server](#))

What <Name> does

- <Section describing problems and the ways that <Name> can solve them

Capabilities

- Data Scientist/Designers - Data scientists/designers can define new types, concepts, and messages for each OnLEP use -case.
- Model Developers/Authors - OnLEP provides lots of model authorship options
 - Define, add/update, activate/deactivate messages, decision rulesets and evaluation models
 - Create, activate and deactivate User-Defined Functions (UDFs) and PMML models for your use-cases.
- Admin - OnLEP provides tools for the platform admin to deploy, configure, monitor, and manage the platform on a cluster of nodes

Ease of Use

- Specify incoming messages by <location, header, xxx>
- Intuitive interface where you can define your own models
- Easy to use rule set definition

<GRAPHIC HERE: useful code snippet>

Performance

<GRAPHIC HERE: useful code snippet>

Fault Tolerance: exactly-once state out of the box.
Speed: 1 Billion transactions per XX (benchmarks here)

Easy to Deploy

Runs on existing Hadoop clusters. Can read data from Kafka and simple file.

BITS - include/delete TBD:

1. <Name> has been designed to be fully customizable and extensible, allowing custom types, concepts, and messages
2. ... is a Big Data and Business Intelligence OLTP Platform. Based on Hadoop, Cassandra,
3. Buzz words to include:
 - a. Big Data, structured and unstructured
 - b. Predictive Analytics
 - c. Data Mining
4. Benefits:
 - a. Ridiculously shorter processing times for terabytes of transactions. Terabyte, Petabyte, Exabyte, Zettabyte, Yottabyte
 - b. Allows powerful strategic insights
 - c. Allows transformative expert system processes