

Status: In Test for v1.0

Wednesday, December 18, 2014, 12:14:48 AM PST

Overview

This guide will help you get the LigaDATA OnLEP engine (Online Learning Engine Platform) and Metadata.API installed and ready to use.

The OnLEP engine processes incoming streaming data, transforms it into messages or containers, and processes it according to the ruleset you have supplied in your models. It then produces decision or alert data that can be acted upon programmatically.

The Metadata.API defines the CRUD operations (create, read, update, delete) for metadata objects supported by this system. The metadata objects include Types, Functions, Concepts, Derived Concepts, MessageDefinitions, and Model Definitions. All functions take String values as input in XML or JSON Format and return a JSON string of ApiResult object.

After you have installed Scala, the JDK, and the RTD engine binary, you can:

- Write your own functions
- Use the JDK to compile those functions

Testers - Is JDK right? Is there something else we should add to this list – one big other thing?

Prerequisites:

- Approximately 300 MB free
- Access to: https://github.com/ligaDATA/RTD_ReadOnly OR the RTD install package

Windows users:

- 7-zip installed (<http://www.7-zip.org/>), choose for 32bit or 64bit.

Note: for OpenSource, we should use zip or tar.gz

- Scala installed, v2.10.4 and install sbt: <http://www.scala-sbt.org/download.html>
- Java installed (Java runtime version 1.7 or later, which can be downloaded [here](#).)
- Install the Java Development Kit (JDK 1.7.1 or greater)
- Download and install Kafka: <http://kafka.apache.org/>
- Download and install Zookeeper: <http://zookeeper.apache.org/releases.html#download>
- The definition of the data models and rule sets used in your business that need to be created as models.

Assumptions:

One instance only of each of the following is running:


- Zookeeper Service
- Kafka
- Cassandra
- HBase

NOTE: Cassandra and HBase are optional - you will need some input/output to the engine.

Linux users: It is suggested that Linux users run the services/scripts under screens (http://www.howtoforge.com/linux_screen), so that if the connection is lost, they will still be running. This will likely become a service in a later release.)

To Test OnLEP on a single node, install the following:

1. From GitHub, clone: https://github.com/ligaDATA/RTD_ReadOnly
 - a. Navigate to the folder:
`<localCloneDirectory>/RTD_ReadOnly/trunk/SampleApplication/Medical/SampleData`
 - b. Confirm that you have the following files:



Linux & Mac users:

2. Run engine install script:

```
<localCloneDirectory>/RTD_ReadOnly/trunk/SampleApplication/Medical/bin/installOnLEP_Medical.sh
```

```
</path/to/install/location> <localCloneDirectory>/RTD_ReadOnly/trunk
```

Example: `<localCloneDirectory>/ligadata/trunk/SampleApplication/Medical/bin/installOnLEP_Medical.sh /tmp/OnLEPInstall $HOME/ligadata/trunk`

Note to testers – the following section is not right – you may have to go back & forth to see what the Config files are doing to figure out what should be happening here:

3. Modify the following two Config files to match your local environment.
 - COPD.cfg – may need to modify **nodeId**. Keep as is for single node test.
 - EngineConfV1.json – may need to modify the Metadata database properties. *This file is used by item 33 in the Metadata.API UI.*
4. Change the local Config file for the Metadata.API for your scenario:
It should be found at:
`<localCloneDirectory>/RTD/trunk/SampleApplication/Medical/Configs/MetadataAPIConfig.properties`

Testers – again I think something is missing here, no idea what.

```
Start MetadataAPI test harness using fat-jar java -jar:  
<localCloneDirectory>/MetadataAPI/target/scala-2.10/MetadataAPI-1.0 --config  
<localCloneDirectory>/MetadataAPIconfig.properties
```

<localCloneDirectory>RTD/trunk/MetadataAPI/target/scala-2.10/MetadataAPI-1.0

5. d

You will see the UI shown here. Enter the number of the item you want to create. It is fairly intuitive from there.



Using the Metadata.API , add metadata objects in the following order:

1. Messages
 - hl7.json
 - inpatientclaim.json
 - outpatientclaim.json
 - beneficiary.json
2. Containers
 - idCodeDim.json
 - CoughCodes.json

- EnvCodes.json
 - SmokeCodes.json
 - SputumCodes.json
 - DyspnoeaCodes.json
 - hcpcsCodes.json
 - icd9DiagnosisCodes.json
 - icd9ProcedureCodes.json
 - dgRelGrpCodes.json
 - IneProcIndTable.json
 - provNumTable.json
3. Models
 - COPDv1.xml
 4. Prepare the kvstores using the new metadata-generated jars:
 5. Run trunk/SampleApplication/Medical/bin/PrepKvStores.sh <installPath> <sourcePath> <engineConfigFile>

Windows:

Note: the script for KVStore initialization will be included in a future release.

Make sure Scala and sbt paths are in the PATH environmental variable.

Example: yours may look different:

```
C:\ProgramData\Oracle\Java\javapath;C:\Program Files\Common Files\Microsoft
Shared\Windows Live;C:\Program Files (x86)\Common Files\Microsoft Shared\Windows
Live;%SystemRoot%\system32;%SystemRoot%;%SystemRoot%\System32\Wbem;%SYSTEMROOT
%\System32\WindowsPowerShell\v1.0\;C:\Program Files (x86)\Windows Live\Shared;;C:\Program
Files (x86)\scala\bin;C:\Program Files (x86)\sbt\bin
```

Install the demo you wish to use:

For instance, execute 'installOnLEP_Medical.bat'

Note: Compression to .gz is commented out in the bat file and needs to be done manually, then copied over manually.

Running OnLEP and processing data

1. Run the engine:
 - a. Execute:


```
/path/to/install/loc/OnLEPManager-1.0 --config </path/to/config/file>
```

Example: /tmp/OnLEPInstall/OnLEPManager-1.0 --config /tmp/OnLEPInstall/COPD.cfg

- b. You can see the input file here:

`/path/to/install/location/msgdata/messages_new_format_all.csv.gz`

Then examine the following files to see the output.

`/path/to/install/location/logs/output1.txt.gz`

Quit the engine by typing "quit". Ctrl-C should also work.

Multi-Node Account Setup Instructions

Each OnLEP cluster installed on a network is ideally run by a distinct user account. These instructions provide the basics for setting up such a user so that it can easily manage the cluster administration.

A user name and group should be selected to be the owner and user that runs the cluster. It can be any useful and meaningful name not already in use. In the examples below, the name *onlep* is used for both the account and group name. Two nodes named *ls19.dc.npario.com* and *ls20.dc.npario.com* are used to illustrate the setup on two machines.

Admin setup items – verify or do the following:

1. If the *onlep* group does not exist on the prospective cluster nodes, add it to each node with the following command:

```
sudo groupadd onlep
```

2. Establish a user account on each prospective cluster node making the new user use the onlep group as its default login group. Be sure to establish a home directory for the user. Example:

```
sudo useradd --gid onlep --create-home onlep
```

3. Set the password for the onlep cluster login to something using sudo

```
sudo passwd onlep
```

This assumes that you can get to these machine from another account and have sudoer privileges.

4. If the onlep release is staged from source on one of the machines in the new cluster or cluster's network, make sure that the staging user (or the developer) account is also a member of the *onlep* group. This can be done with this command:

```
sudo usermod -G onlep releaseAdmin
```

where the user account is releaseAdmin. This really is only needed on this one machine, not every node in the cluster, although THAT would not necessarily be a bad idea. This is security policy dependent.

5. Passphrase-less access to/from all participating nodes in the *onlep* cluster required for proper operation. Each node could have its own public and private key if your security policy requires it. However, it is far easier to login to one of the nodes as the onlep user and set up the .ssh keys once,

then propagate the entire directory to the other nodes in the cluster. Assuming the accounts are new, there is really no harm in replacing the entire .ssh directory. If you have existing keys there, however, don't replace. Instead append.

To append to the .ssh keys

1. generate a public private key

```
ssh-keygen
```

It will prompt for a key type. Take the default (RSA 2048) by hitting return. Then it will ask for a passphrase; hit return for "passphraseless" behavior. It will prompt again to type the same passphrase; once again hit return. This will generate a public and private key for the current user account.

2. Install the public key on the current machine.

To install the new public key in the current directory, you can cat the public key into a file called .ssh/authorized_keys or perhaps a little easier use this approach, as it will properly set the access permissions for the file it will create. If the current machine was named ls19.dc.npario.com and the onlep account being setup is onlep, do the following:

```
ssh-copy-id onlep@ls19.dc.npario.com
```

You will have to type the password this one time to establish the authorized key file, but after that you should be able to access without giving a passphrase.

3. Verify that the \$HOME/.ssh permissions is strict. There should be no permitted access to the .ssh directory except by the owner.
4. As the onlep account, Zip/Bzip/Gzip the .ssh folder in its entirety and leave it in \$HOME. For example:

```
tar cvjf onlepNodeSSH.bz2 .ssh
```

This creates an archive that will be copied around to the other nodes that are to participate in the cluster.

5. As the onlep account, copy the archive to every node in the cluster and unzip it. For example, if ls20.dc.npario.com were another node in the cluster to be setup, do the following:

```
scp onlepNodeSSH.bz2 ls20.dc.npario.com:
```

This will leave the tarball in the onlep user's ls20 home folder. Do this same action for all the nodes to participate or might participate in the cluster being set up.

6. ssh to each node as the onlep user and decompress the file copied in the previous step. For example,

```
tar xvjf onlepNodeSSH.bz2
```

This will update the .ssh folder with the authorized_keys as well as the public and private keys that were generated in part a) above. To verify that it worked, try to ssh to the original machine where the keys were generated. For example,

```
ssh ls19.dc.npario.com
```

If successful, then you should login immediately (not requiring a passphrase from the console).

7. Check further. A good check would be to login to one of the cluster nodes to participate in the onlep cluster being formed AS the onlep account and try to log in to every other machine in the cluster. These logins should succeed without a passphrase being required.

Should there be problems, there are many resources on the web that describe the setup of a passphrase-less account. In particular, the access permissions for the account and its .ssh need to be strict (no group or global permissions should be allowed). Ssh will complain if these requirements are not met.

8. Finally, remove the tarball that has been passed around to all of the nodes from each onlep user account on each node.

More information. The instructions above are indicative, not authoritative, of what must be done to establish password-less access. If you experience problems, consult the documentation for your Linux enlistment for special requirements that your environment may require.

Installing OnLEP from your Git Repository

OnLEP is typically installed from the trunk directory of an “install” or “release” machine, building the OnLEP installation directory on the local system, creating a tarball of that directory, sending it to the nodes that are part of the OnLEP cluster, untar and decompress the tarball there and move the resulting directory into place.

Several scripts are part of the repository that are used to accomplish this. These scripts must be moved onto your PATH and given “execute” permissions. They are:

1. trunk/SampleApplication/Medical/bin/clusterInstallOnLEP.sh

2. `trunk/SampleApplication/Medical/bin/installOnLEP_Medical.sh`
3. `trunk/SampleApplication/Medical/bin/nodeInfoExtract.sh`
4. `trunk/Pmml/Scripts/sbtProjDependencies/src/main/scala/sbtProjDependencies.scala`

The *clusterInstallOnLEP.sh script* controls the installation and uses the other scripts mentioned. The *installOnLEP_Medical.sh* will build the OnLEP install directory. The *nodeInfoExtract.sh* is used to determine which cluster nodes are to receive the build. The *sbtProjDependencies* is a useful tool used by a number of scripts that are part of OnLEP. In this case, it is used to build the classpath for the *nodeInfoExtract* application.

Invoke the installer

Use the following command by invoking it from the trunk directory of the local Git repository which is to be installed. Run it from the account that is to be used to run/manage the OnLEP cluster and refer to directories that are write-able by that account.

```
Cluster InstallOnLEP.sh
--MetadataAPIConfig SampleApplication/Medical/Configs/MetadataAPIConfig.properties
--NodeConfigPath SampleApplication/Medical/Configs/Engine2BoxConfigV1.json
```

Currently two configuration files are required, one for the MetadataAPI and the other that describes at least the cluster for the OnLEPManager. The cluster description is the key information as to what sort of distribution and what machines and nodes are involved in the installation. Note that multiple nodes can if desired be installed on the same physical computer.

For more information about what these configuration files contain, see <some reference to other documentation describing the configuration>.